

实验四

PB17081504 廖洲洲

实验题目

利用MPI实现并行排序算法

实验环境

操作系统: Win 10

IDE: Microsoft Visual Studio 2015

编译器: cl.exe

硬件配置: CPU: Intel Core i7-8550U; CPU核心数:4; 内存:8G

算法设计与分析

PSRS算法

Phase I: Initialization

Start up p processors, let the root processor 0 get data of size n .

Phase II: Scatter data, local sort and regular samples collected

Scatter the data values to the p processors. Each processors sorts its local data set, roughly of size n/p , using *quicksort*. Each processor chooses p sample points, in a very regular manner, from its locally sorted data.

Phase III: Gather and merge samples, choose and broadcast $p-1$ pivots

The root processor 0 gathers the p sets of p sample points. It is important to realize that each set of these p sample points is sorted. These p sets are sorted using *multimerge*. From these p^2 sorted points, $p-1$ pivot values are regularly chosen and are *broadcast* to the other $p-1$ processors.

Phase IV: Local data is partitioned

Each of the p processors *partitions* its local sorted data, roughly of size n/p , into p classes using the $p-1$ pivot values.

Phase V: All i^{th} classes are gathered and merged

Processor i gathers the i^{th} class of data from every other processor. Each of these classes is sorted using *multimerge*.

Phase VI: Root processor collects all the data

The root processor *gathers* all the data and assembles the sorted list of n values.

快速排序的平均复杂度为 $O(n\log n)$,使用的归并排序花费的时间为 $O(n)$,因此单个进程的所需的时间复杂度为 $O(n\log n)$ 。空间复杂度为 $O(n)$

核心代码

```
startTime = MPI_Wtime();
//均匀划分, 将待处理数据均分为size份, 进程rank处理下标rank*partition~
(rank+1)*partition-1段数组
partition = N / size;
int start = rank*partition;
//局部排序, 使用快速排序对本地数据进行排序
quickSort(num + rank*partition, partition);
//选取样本, 从排序好的子序列中选取size个样本, 并发送给0号进程, 0号进程使用Gather收集数据
double *cbuffer = (double *)malloc(sizeof(double)*size);
int space = partition / size;
for (i = 0; i < size; i++) {
    cbuffer[i] = (num + rank*partition)[i*space];
}
double *pivotbuffer = (double *)malloc(sizeof(double)*size*size);
MPI_Barrier(MPI_COMM_WORLD);
MPI_Gather(cbuffer, size, MPI_DOUBLE, pivotbuffer, size, MPI_DOUBLE, 0,
MPI_COMM_WORLD);
//样本排序, 对选取的size*size个样本进行多次归并排序
if (rank == 0) {
    for (i = 1; i < size; i++) {
        merge(pivotbuffer, 0, i*size-1, (i + 1)*size-1);
    }
    //选择size-1个主元
    for (i = 1; i < size; i++) {
        cbuffer[i] = pivotbuffer[i*size - 1];
    }
}
//使用广播将主元播送给其他进程
MPI_Bcast(cbuffer, size, MPI_DOUBLE, 0, MPI_COMM_WORLD);
//主元划分, 根据主元将本地数据划分为size段
int *classStart = (int *)malloc(sizeof(int)*size);
int *classListLength = (int *)malloc(sizeof(int)*size);
for (i = 0; i < size; i++) {
    classStart[i] = 0;
    classLength[i] = 0;
}
int index = 0;
j = 1;
for (i = 0; i < partition; i++) {
    if (j<size&& num[start + i] > cbuffer[j]) {
        classStart[j] = i;
        j++;
        index++;
    }
    classLength[index]++;
}
//接收各段长度
int *recvLength = (int *)malloc(sizeof(int)*size);
int *recvStart = (int *)malloc(sizeof(int)*size);
for (i = 0; i < size; i++) {
    MPI_Gather(classLength + i, 1, MPI_INT, recvLength, 1, MPI_INT, i,
MPI_COMM_WORLD);
}
recvStart[0] = 0;
```

```

//计算接收的各段的起始地址
for (i = 1; i < size; i++) {
    recvStart[i] = recvStart[i - 1] + recvLength[i-1];
}
//交换，各处理器将其有序段按段号交换到对应的处理器中
double *recvBuffer = (double *)malloc(sizeof(double)*N);
for (i = 0; i < size; i++) {
    MPI_Gatherv(num + start + classStart[i], classLength[i], MPI_DOUBLE,
recvBuffer, recvLength, recvStart, MPI_DOUBLE, i, MPI_COMM_WORLD);
}
MPI_Barrier(MPI_COMM_WORLD);
//归并排序，各处理器将收到的元素进行归并排序
for (i = 1; i < size; i++) {
    merge(recvBuffer, 0, recvStart[i] - 1, recvStart[i]+recvLength[i]-1);
}
//收集，0号进程收集各进程处理的数据
int len = 0;
for (i = 0; i < size; i++) {
    len += recvLength[i];
}
//将各个进程归并处理的数据长度发给0号进程
MPI_Gather(&len, 1, MPI_INT, recvLength, 1, MPI_INT, 0, MPI_COMM_WORLD);
recvStart[0] = 0;
for (i = 1; i < size; i++) {
    recvStart[i] = recvStart[i - 1] + recvLength[i-1];
}
MPI_Gatherv(recvBuffer, len, MPI_DOUBLE, num, recvLength, recvStart,
MPI_DOUBLE, 0, MPI_COMM_WORLD);
endTime = MPI_Wtime();

```

实验结果

运行时间

规模\进程数	1	2	4	8
10000	0.001100	0.001345	0.001975	0.004912
100000	0.013461	0.009364	0.007234	0.010221
1000000	0.215377	0.102547	0.067189	0.063981
10000000	19.275609	5.368064	1.799027	0.974133

加速比

规模\进程数	1	2	4	8
10000	1	0.817844	0.556962	0.223941
100000	1	1.437527	1.860796	1.316994
1000000	1	2.100276	3.20554	3.366265
10000000	1	3.590793	10.71446	19.78745

分析与总结

- PSRS算法使用了多次和多种通信，让我对MPI的各种通信方法有了更加深刻的理解。
- PSRS算法多个进程保留了所有数据并且还设置了接收buffer，故空间的使用会较串行算法更高些
- 当数据规模较小时，并行的加速效果并不明显，甚至进程间的大量通信会导致程序运行的时间更长
- 当数据规模增大，并行的加速效果越来越显著，在开启8个进程的情况下对1千万个数据进行排序，时间花费竟然小于1s，加速比达到了近20，效果惊人。

实验截图

数据规模N=64

- 进程数=1

```
E:\Visual Studio Projects\MPI_PSRS\x64\Debug>mpiexec -n 1 MPI_PSRS.exe
Before sort:
41.00 57.00 1.00 16.00 8.00 32.00 76.00 5.00 55.00 83.00 31.00 25.00 79.00 23.00 89.00 96.00 22.00 69.00 4.00 61.00 70.00 66.00 77.00
74.00 49.00 19.00 10.00 37.00 86.00 6.00 23.00 59.00 80.00 94.00 82.00 1.00 45.00 14.00 29.00 73.00 63.00 14.00 90.00 50.00 22.00 56.00
41.00 67.00 71.00 73.00 99.00 66.00 60.00 53.00 39.00 56.00 2.00 62.00 86.00 52.00 39.00 72.00 87.00 65.00
After sort:
1.00 1.00 2.00 4.00 5.00 6.00 8.00 10.00 14.00 14.00 16.00 19.00 22.00 22.00 23.00 23.00 25.00 29.00 31.00 32.00 37.00 39.00 39.00 41.
00 41.00 45.00 49.00 50.00 52.00 53.00 55.00 56.00 56.00 57.00 59.00 60.00 61.00 62.00 63.00 65.00 66.00 66.00 67.00 69.00 70.00 71.00
72.00 73.00 73.00 74.00 76.00 77.00 79.00 80.00 82.00 83.00 86.00 86.00 87.00 89.00 90.00 94.00 96.00 99.00
Time:0.000026
```

- 进程数=2

```
E:\Visual Studio Projects\MPI_PSRS\x64\Debug>mpiexec -n 2 MPI_PSRS.exe
Before sort:
41.00 11.00 54.00 5.00 32.00 7.00 94.00 99.00 14.00 52.00 21.00 28.00 33.00 84.00 46.00 13.00 4.00 62.00 33.00 12.00 71.00 53.00 31.00
67.00 74.00 95.00 6.00 93.00 88.00 36.00 34.00 43.00 58.00 99.00 26.00 8.00 90.00 75.00 51.00 7.00 5.00 3.00 8.00 24.00 45.00 38.00 4
7.00 25.00 22.00 57.00 57.00 42.00 90.00 4.00 34.00 41.00 49.00 42.00 55.00 54.00 51.00 93.00 54.00 17.00
After sort:
3.00 4.00 4.00 5.00 5.00 6.00 7.00 7.00 8.00 8.00 11.00 12.00 13.00 14.00 17.00 21.00 22.00 24.00 25.00 26.00 28.00 31.00 32.00 33.00
33.00 34.00 34.00 36.00 38.00 41.00 41.00 42.00 42.00 43.00 45.00 46.00 47.00 49.00 51.00 51.00 52.00 53.00 54.00 54.00 54.00 55.00 57
.00 57.00 58.00 62.00 67.00 71.00 74.00 75.00 84.00 88.00 90.00 90.00 93.00 93.00 94.00 95.00 99.00 99.00
Time:0.000390
```

- 进程数=4

```
E:\Visual Studio Projects\MPI_PSRS\x64\Debug>mpiexec -n 4 MPI_PSRS.exe
Before sort:
41.00 74.00 91.00 93.00 83.00 3.00 70.00 28.00 17.00 2.00 17.00 29.00 74.00 88.00 88.00 36.00 59.00 12.00 77.00 86.00 50.00 31.00 98.00
38.00 30.00 65.00 26.00 96.00 32.00 34.00 7.00 59.00 98.00 27.00 59.00 75.00 40.00 9.00 62.00 27.00 84.00 16.00 84.00 28.00 5.00 50.0
0 18.00 54.00 27.00 8.00 95.00 99.00 88.00 74.00 28.00 70.00 61.00 5.00 70.00 3.00 12.00 22.00 38.00 35.00
After sort:
2.00 3.00 3.00 5.00 5.00 7.00 8.00 9.00 12.00 12.00 16.00 17.00 17.00 18.00 22.00 26.00 27.00 27.00 27.00 27.00 28.00 28.00 28.00 29.00 30.0
0 31.00 32.00 34.00 35.00 36.00 38.00 38.00 40.00 41.00 50.00 50.00 54.00 59.00 59.00 59.00 61.00 62.00 65.00 70.00 70.00 70.00 74.00
74.00 74.00 75.00 77.00 83.00 84.00 84.00 86.00 88.00 88.00 88.00 91.00 93.00 95.00 96.00 98.00 98.00 98.00 99.00
Time:0.001585
```

- 进程数=8

```
E:\Visual Studio Projects\MPI_PSRS\x64\Debug>mpiexec -n 8 MPI_PSRS.exe
Before sort:
41.00 97.00 26.00 88.00 81.00 31.00 51.00 58.00 75.00 45.00 9.00 31.00 57.00 48.00 75.00 41.00 14.00 81.00 84.00 50.00 12.00 29.00 68.0
0 45.00 77.00 6.00 61.00 51.00 11.00 16.00 81.00 84.00 6.00 18.00 87.00 98.00 59.00 2.00 28.00 51.00 36.00 98.00 63.00 49.00 98.00 27.
00 88.00 97.00 26.00 17.00 34.00 97.00 70.00 91.00 76.00 40.00 52.00 53.00 31.00 31.00 44.00 78.00 64.00 78.00
After sort:
2.00 6.00 6.00 9.00 11.00 12.00 14.00 16.00 17.00 18.00 26.00 26.00 27.00 28.00 29.00 31.00 31.00 31.00 31.00 51.00 34.00 36.00 40.00
41.00 41.00 44.00 45.00 51.00 61.00 45.00 48.00 49.00 50.00 51.00 52.00 59.00 70.00 77.00 53.00 57.00 58.00 63.00 68.00 76.00 81.00 87.
00 64.00 75.00 75.00 78.00 78.00 81.00 81.00 84.00 88.00 91.00 98.00 84.00 88.00 97.00 97.00 97.00 98.00 98.00
Time:0.003139
```